

# Homework 4

1. [10 points] Design a data structure that has the following properties (assume  $n$  elements in the data structure, and that the data structure properties need to be preserved at the end of each operation):

- a. • Find median takes  $O(1)$  time
- b. • Insert takes  $O(\log n)$  time

Do the following:

1. Describe how your data structure will work.
2. Give algorithms that implement the Find-Median() and Insert() functions.

**Solution.** We use the  $dn/2e$  smallest elements to build a max-heap and use the remaining  $bn/2c$  elements to build a min-heap. The median will be at the root of the max-heap and hence accessible in time  $O(1)$  (we assume the case of even  $n$ , median is  $n/2$ -th element when elements are sorted in increasing order).

Insert() algorithm: For a new element  $x$

- Initialize len of maxheap (maxlen) and len of minheap (minlen) to 0.
- Compare  $x$  to the current root of max-heap.
- If  $x < \text{median}$ , we insert  $x$  into the max-heap. Maintain length of maxheap say maxlen, and every time you insert element into maxheap increase maxlen by 1. Otherwise, we insert  $x$  into the min-heap, and increase length of minheap say minlen by 1. This takes  $O(\log n)$  time in the worst case.
- If  $\text{size}(\text{maxHeap}) > \text{size}(\text{minHeap}) + 1$ , then we call Extract-Max() on max heap, and decrease maxlen by 1 of and insert the extracted value into the min-heap and increase minlen by 1. This takes  $O(\log n)$  time in the worst case.
- Also, if  $\text{size}(\text{minHeap}) > \text{size}(\text{maxHeap})$ , we call Extract-Min() on min heap, decrease minlen by 1 and insert the extracted value into the max heap and increase maxlen by 1. This takes  $O(\log n)$  time in the worst case.

Find-Median() algorithm:

- If  $(\text{maxlen} + \text{minlen})$  is even: return  $(\text{sum of roots of max heap and min heap})/2$  as median
- Else if  $\text{maxlen} > \text{minlen}$ : return root of max heap as median
- Else: return root of min heap as median

Rubric

- 3 pt: Using max heap and min heap to store first half and second half of elements.

- 3 pt: Comparing element to the root and proper if conditions for inserting in appropriate heap.
- 2 pt: Proper if conditions for returning the median.
- 2 pt: Correct Time Complexity

2. [20 points] A network of  $n$  servers under your supervision is modeled as an undirected graph  $G = (V, E)$  where a vertex in the graph corresponds to a server in the network and an edge models a link between the two servers corresponding to its incident vertices. Assume  $G$  is connected. Each edge is labeled with a positive integer that represents the cost of maintaining the link it models. Further, there is one server (call its corresponding vertex as  $S$ ) that is not reliable and likely to fail. Due to a budget cut, you decide to remove a subset of the links while still ensuring connectivity. That is, you decide to remove a subset of  $E$  so that the remaining graph is a spanning tree. Further, to ensure that the failure of  $S$  does not affect the rest of the network, you also require that  $S$  is connected to exactly one other vertex in the remaining graph. Design an algorithm that given  $G$  and the edge costs efficiently decides if it is possible to remove a subset of  $E$ , such that the remaining graph is a spanning tree where  $S$  is connected to exactly one other vertex and (if possible) finds a solution that minimizes the sum of maintenance costs of the remaining edges.

### Solution

First we need to check the possibility of node  $S$  having only one neighbor in a spanning tree of the underlying graph. The best way to check this is to remove node  $S$  and all its adjacent edges to form a graph  $G'$ . If  $G'$  is a connected graph, then we can claim it is possible to have a spanning tree where node  $S$  has only one neighbor. To check the connectivity of  $G'$ , the simplest way is to run a DFS or BFS algorithm on  $G'$ .

Considering that  $G'$  is connected, we need to find the spanning tree that minimizes the maintenance cost. Therefore we need to find the MST with the additional constraint that  $S$  should be a leaf. Therefore, we remove  $S$  and all its adjacent edges to form  $G'$ . We run Prime's (or any other MST algorithm) to find the MST of  $G'$ . Among all edges adjacent to  $S$ , we find the one with the minimum maintenance cost and connect  $S$  to the MST using this edge. The resulting graph will still be a spanning tree and in this spanning tree  $S$  will be a leaf.

### Rubric

- 5 pt: provides a way to check whether such an MST is possible
- 10 pt: finds MST for a graph after removing all edges to/from  $S$
- 5 pt: states least cost edge from  $S$  is added to MST.

3. [15 points] Prove or disprove the following:

- $T$  is a spanning tree on an undirected graph  $G = (V, E)$ . Edge costs in  $G$  are NOT guaranteed to be unique. If every edge in  $T$  belongs to SOME minimum cost spanning trees in  $G$ , then  $T$  is itself a minimum cost spanning tree.
- Consider two positively weighted graphs  $G = (V, E, w)$  and  $G' = (V, E, w')$  with the same vertices  $V$  and edges  $E$  such that, for any edge  $e$  in  $E$ , we have  $w'(e) = w(e)^2$ . For any two vertices  $u, v$  in  $V$ , any shortest path between  $u$  and  $v$  in  $G'$  is also a shortest path in  $G$ .

### Solution

1. False. Counter example:

$ab = 2, bc = 2, ac = 1$

$T = ab, bc$

$MST1 = ab, ac$

$MST2 = ac, bc$

$ab$  is in  $MST1$

$bc$  is in  $MST2$

However,  $T$  is not a MST

### Rubric

- 3 pt: Correct T/F claim
- 5 pt: Provides a correct counterexample as explanation

2. False. Assume we have two paths in  $G$ , one with weights 2 and 2 and another one with weight 3. The first one is shorter in  $G'$  while the second one is shorter in  $G$ .

### Rubric

- 3 pt: Correct T/F claim
- 4 pt: Provides a correct counterexample as explanation

4. [15 points] A new startup FastRoute wants to route information along a path in a communication network, represented as a graph. Each vertex represents a router and each edge a wire between routers. The wires are weighted by the maximum bandwidth they can support. FastRoute comes to you and asks you to develop an algorithm to find the

path with maximum bandwidth from any source  $s$  to any destination  $t$ . As you would expect, the bandwidth of a path is the minimum of the bandwidths of the edges on that path; the minimum edge is the bottleneck. Explain how to modify Dijkstra's algorithm to do this.

### Solution

Data structure change: we'll use a max heap instead of a min heap used in Dijkstra's.

Initialization of the heap: Initially all nodes will have a distance (bandwidth) of zero from  $s$ , except the starting point  $s$  which will have a bandwidth of  $\infty$  to itself.

Change in relaxation step. Based on the definition of a path's bandwidth, the bandwidth of a path from  $s$  to  $u$  through  $u$ 's neighbor  $v$  will be  $d(u) = \min(d(v), \text{weight}(v, u))$ , because the bandwidth of a path is equal to the bandwidth of its lowest bandwidth edge. Therefore, in the relaxation step, we will be replacing:

$$d(u) = \min(d(u), d(v) + \text{weight}(v, u))$$

with

$$d(u) = \max(d(u), \min(d(v), \text{weight}(v, u)))$$

### Rubric

- 2 points for max heap
- 3 points for initialization
- 10 points for relaxation modification

5. [10 points] There is a stream of integers that comes continuously to a small server. The job of the server is to keep track of  $k$  largest numbers that it has seen so far. The server has the following restrictions:
- It can process only one number from the stream at a time, which means it takes a number from the stream, processes it, finishes with that number and takes the next number from the stream. It cannot take more than one number from the stream at a time due to memory restriction.
  - It has enough memory to store up to  $k$  integers in a simple data structure (e.g. an array), and some extra memory for computation (like comparison, etc.).
  - The time complexity for processing one number must be better than  $O(k)$ . Anything that is  $O(k)$  or worse is not acceptable. Design an algorithm on the server to perform its job with the requirements listed above.

Solution. Use a binary min-heap on the server.

1. Do not wait until  $k$  numbers have arrived at the server to build the heap, otherwise you would incur a time complexity of  $O(k)$ . Instead, build the heap on-the-fly, i.e. as soon as a number arrives, if the heap is not full, insert the number into the heap and execute `Heapify()`. The first  $k$  numbers are obviously the  $k$  largest numbers that the server has seen.
2. When a new number  $x$  arrives and the heap is full, compare  $x$  to the minimum number  $r$  in the heap located at the root, which can be done in  $O(1)$  time. If  $x \leq r$ , ignore  $x$ . Otherwise, run `Extract-min()` and insert the new number  $x$  into the heap and call `Heapify()` to maintain the structure of the heap.
3. Both `Extract-min()` and `Heapify()` can be done in  $O(\log k)$  time. Hence, the overall complexity is  $O(\log k)$ .

#### Rubric

- 2 points - states min-heap as the desired data structure
- 3 points - correct method for inserting when heap is not full
- 4 points - correct method for inserting when heap is full
- 1 point - correct time complexity

6. [15 points] Consider a directed, weighted graph  $G$  where all edge weights are positive. You have one Star, which allows you to change the weight of any one edge to zero. In other words, you may change the weight of any one edge to zero. Propose an efficient method based on Dijkstra's algorithm to find a lowest-cost path from node  $s$  to node  $t$ , given that you may set one edge weight to zero.

Note: you will receive 10 points if your algorithm is efficient. This means your method must do better than the naive solution, where the weight of each node is set to 0 per time and the Dijkstra's algorithm is applied every time for lowest-cost path searching. You will receive full points (15 points) if your algorithm has the same run time complexity as Dijkstra's algorithm.

Solution. Use Dijkstra's algorithm to find the shortest paths from  $s$  to all other vertices. Reverse all edges and use Dijkstra to find the shortest paths from all vertices to  $t$ . Denote the shortest path from  $u$  to  $v$  by  $u \rightarrow v$ , and its length by  $\delta(u, v)$ . Now, try setting each edge to zero. For each edge  $(u, v) \in E$ , consider the path  $s \rightarrow u \rightarrow v \rightarrow t$ . If we set  $w(u, v)$  to zero, the path length is  $\delta(s, u) + \delta(v, t)$ . Find the edge for which this length is minimized and set it to zero; the corresponding path  $s \rightarrow u \rightarrow v \rightarrow t$  is the desired path. The algorithm requires two invocations of Dijkstra, and an additional  $O(E)$  time to iterate through the edges and find the optimal edge to

take for free. Thus the total running time is the same as that of Dijkstra:

- If we implement Dijkstra's algorithm with Fibonacci heap
  - time complexity of both the Dijkstra's algorithm and the proposed method is  $O(E + V \log V)$ .
  - the corresponding time complexity of naive is  $O(E(E + V \log V))$
- If we implement Dijkstra's algorithm with a binary heap
  - the complexity of Dijkstra's algorithm is  $O((E + V) \log V)$ , so that the proposed method is  $O(2(E + V) \log V + E)$ , since  $E = O(E \log V)$ , then it has same time complexity as Dijkstra's algorithm, i.e.,  $O((E + V) \log V)$ .
  - the corresponding time complexity of naive is  $O(E(E + V) \log V)$

Rubrics:

- 12 pt if the approach is of the same complexity as Dijkstra's algorithm
- 4 pt if the approach is more efficient than naive but not have the same complexity as Dijkstra's algorithm
- 0 pt if the naive approach or any other approach with larger complexity than naive is proposed
- 3 pt for time complexity analysis.

7. [10 points] When constructing a binomial heap of size  $n$  using  $n$  insert operations, what is the amortized cost of the insert operation? Find using the accounting method.

Solution:

We make sure that at any point, each tree has a credit of 1. Insert involves creating a new heap with just one element (which is constant time) and then merging if there are any other heaps with only 1 element and so on.

Let's say inserting a tree with a single new element has an amortized cost of 2 and an actual cost of 1. This means that for every insert we get a credit of 1. This credit is used to merge two trees. We use the credit stored in the tree which is merged as the left child of the other root. Thus the insert operation ends up having an amortized cost of  $\Theta(1)$ . It is conceptually similar to a binary counter, where each bit flip from 0 to 1 is assigned an amortized cost of 2 while having an actual cost of 1, and the credit created here is used to flip the bit from 1 to 0. In the case of a binomial heap, we have  $\log n$  bits and each tree (the bit which is 1) has 1 credit. One of the credits is used to merge while the credit from the tree it is being merged with remains.

Rubrics:

- 2 pt - mentions each tree must have 1 credit associated with it
- 2 pt - Correctly assigns amortized cost of 2 to each insert
- 2 pt - Correctly assigns actual cost of 1 to each insert
- 2 pt - correctly says joining of two trees maintains the credit with the combined tree
- 2 pt - states the amortized cost is  $\Theta(1)$